

Shiv Nadar University Chennai

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

Experiment 1 Implementation of a Single Layer Perceptron for Binary Classification

1. Objective

The objective of this experiment is to understand the concept of an artificial neuron and implement a Single Layer Perceptron from scratch for binary classification. Students will learn the perceptron learning algorithm, understand the importance of activation functions, visualize the learning process, and evaluate the classifier using a real-world dataset.

2. Tasks to be Performed

1. Download and understand the dataset.
2. Perform Exploratory Data Analysis (EDA).
3. Preprocess the dataset.
4. Implement a Single Layer Perceptron from scratch.
5. Train the model using the perceptron learning rule.
6. Evaluate the classifier using standard performance metrics.
7. Analyze the effect of different learning rates.
8. Compare the implementation with Scikit-learn's Perceptron (optional).

3. Background Theory

An Artificial Neuron is the fundamental building block of a neural network. It receives one or more input features, multiplies them by corresponding weights, adds a bias term, and applies an activation function to determine the output.

The Single Layer Perceptron is the earliest supervised learning algorithm capable of solving linearly separable binary classification problems.

Mathematical Formulation

Weighted Sum

$$z = \mathbf{w}^T \mathbf{x} + b$$

where

- $\mathbf{x}$: Input vector
- $\mathbf{w}$: Weight vector
- b : Bias
- z : Linear combination

Prediction

$$\hat{y} = f(z)$$

Weight Update Rule

$$\mathbf{w}_{new} = \mathbf{w}_{old} + \eta(y - \hat{y})\mathbf{x}$$

Bias Update

$$b_{new} = b_{old} + \eta(y - \hat{y})$$

where η denotes the learning rate.

4. Activation Functions

An activation function determines the output of a neuron based on the weighted sum of its inputs. It introduces non-linearity, enabling neural networks to learn complex patterns. Although the perceptron uses only the Step Activation Function, modern deep learning models employ differentiable activation functions for efficient optimization through backpropagation.

Step Activation Function

$$f(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$$

The Step Activation Function produces a binary output and is therefore suitable only for linearly separable binary classification problems.

Activation	Output Range	Typical Applications
Step	$\{0, 1\}$	Classical Perceptron
Sigmoid	$(0, 1)$	Binary Classification Output Layer
Tanh	$(-1, 1)$	Hidden Layers (Traditional Networks)
ReLU	$[0, \infty)$	Hidden Layers in CNNs and MLPs
Leaky ReLU	$(-\infty, \infty)$	Avoiding Dying ReLU
Softmax	$(0, 1)$	Multi-class Classification Output Layer

Note: This experiment uses only the Step Activation Function. The remaining activation functions will be studied in subsequent experiments involving Multilayer Perceptrons and Deep Neural Networks.

5. Dataset Description

Dataset: Banknote Authentication Dataset

Source: UCI Machine Learning Repository

Download Link:

<https://archive.ics.uci.edu/dataset/267/banknote+authentication>

Instances	1372
Features	4 Numerical Features
Classes	2
Missing Values	None
Task	Binary Classification

Feature Description

- Variance
- Skewness
- Curtosis
- Entropy

Target Class

- 0 – Authentic Banknote
- 1 – Forged Banknote

6. Experimental Procedure and Results

Task 1: Dataset Exploration

- Load the dataset.
- Display the first five samples.
- Determine dataset dimensions.
- Identify missing values.
- Display descriptive statistics.

Output

First five samples:					
	Variance	Skewness	Curtosis	Entropy	Class
0	3.62160	8.6661	-2.8073	-0.44699	0
1	4.54590	8.1674	-2.4586	-1.46210	0
2	3.86600	-2.6383	1.9242	0.10645	0
3	3.45660	9.5228	-4.0112	-3.59440	0
4	0.32924	-4.4552	4.5718	-0.98880	0

```
Dataset dimensions (rows, columns): (1372, 5)
```

```
Missing values per column:
```

```
Variance      0  
Skewness      0  
Curtosis      0  
Entropy       0
```

```
Class         0
```

```
dtype: int64
```

```
Descriptive statistics:
```

	Variance	Skewness	Curtosis	Entropy	Class
count	1372.000000	1372.000000	1372.000000	1372.000000	1372.000000
mean	0.433735	1.922353	1.397627	-1.191657	0.444606
std	2.842763	5.869047	4.310030	2.101013	0.497103
min	-7.042100	-13.773100	-5.286100	-8.548200	0.000000
25%	-1.773000	-1.708200	-1.574975	-2.413450	0.000000
50%	0.496180	2.319650	0.616630	-0.586650	0.000000
75%	2.821475	6.814625	3.179250	0.394810	1.000000
max	6.824800	12.951600	17.927400	2.449500	1.000000

```
Class distribution:
```

```
Class
```

```
0      762
```

```
1      610
```

```
Name: count, dtype: int64
```

Task 2: Exploratory Data Analysis

Generate Histograms, Correlation Heatmap, Scatter Plot, and Boxplots.




Figure 1: Feature Histograms




Figure 2: Correlation Heatmap




Figure 3: Scatter Plot: Variance vs Skewness




Figure 4: Boxplots of Features by Class

Task 3: Data Preprocessing

- Normalize all numerical features.
- Split the dataset into Training (80%) and Testing (20%).

Output

```
Training set size: (1097, 4)
Testing set size: (275, 4)
```

Task 4 & 5: Perceptron Implementation and Model Training

Implement from scratch: Weight Initialization, Bias Initialization, Step Activation Function, Forward Propagation, and the Perceptron Learning Rule. Display, per epoch: Epoch Number, Misclassified Samples, Updated Weights, Updated Bias.

Epoch	1		Misclassified:	59		Weights:	[-0.068 -0.0938 -0.0678 -0.011]		Bias:	-0.03
Epoch	2		Misclassified:	25		Weights:	[-0.0775 -0.1042 -0.0811 0.0035]		Bias:	-0.04
Epoch	3		Misclassified:	20		Weights:	[-0.0802 -0.1127 -0.0917 -0.0036]		Bias:	-0.04
Epoch	4		Misclassified:	19		Weights:	[-0.0883 -0.1196 -0.1001 0.0089]		Bias:	-0.05
Epoch	5		Misclassified:	20		Weights:	[-0.0991 -0.1295 -0.1037 0.0102]		Bias:	-0.05
...										
Epoch	46		Misclassified:	15		Weights:	[-0.1942 -0.2323 -0.2175 -0.0208]		Bias:	-0.09
Epoch	47		Misclassified:	16		Weights:	[-0.1987 -0.2279 -0.218 -0.0179]		Bias:	-0.09
Epoch	48		Misclassified:	17		Weights:	[-0.1951 -0.2336 -0.2113 -0.0119]		Bias:	-0.10
Epoch	49		Misclassified:	14		Weights:	[-0.2092 -0.2231 -0.2069 -0.0255]		Bias:	-0.10
Epoch	50		Misclassified:	18		Weights:	[-0.1941 -0.2358 -0.2109 -0.0195]		Bias:	-0.10

Task 6: Model Evaluation

Compute Accuracy, Precision, Recall, F1-score, and the Confusion Matrix.

Output

```
Accuracy : 0.9855
Precision : 0.9683
Recall : 1.0000
F1-score : 0.9839
```

Confusion Matrix:

```
[[149  4]
 [ 0 122]]
```




Figure 5: Confusion Matrix



Figure 6: Training Error vs Epoch




Figure 7: Weight Evolution over Epochs




Figure 8: Bias Evolution over Epochs

Task 7: Effect of Different Learning Rates

Output

Learning Rate = 0.001	Epochs run = 50	Final Misclassified = 18	Test Accuracy = 0.9855
Learning Rate = 0.01	Epochs run = 50	Final Misclassified = 18	Test Accuracy = 0.9855
Learning Rate = 0.1	Epochs run = 50	Final Misclassified = 18	Test Accuracy = 0.9855




Figure 9: Learning Rate Comparison: Convergence Behaviour

Optional: Decision Boundary (Variance vs Skewness)




Figure 10: Decision Boundary: Variance vs Skewness

7. Mandatory Plots Summary

Plot	Purpose	Expected Inference
Feature Histograms	Distribution Analysis	Identify skewness and outliers.
Correlation Heatmap	Feature Relationship	Observe highly correlated features.
Scatter Plot	Class Distribution	Checks if classes are linearly separable.
Training Error vs Epoch	It learns behaviour	Verify convergence of the perceptron.
Weight Evolution	Parametric Analysis	Observe how stable the learned weights are.
Bias Evolution	Parametric Analysis	Does the role of bias term.
Confusion Matrix	Performance Evaluation	Interpret TP, FP, TN and FN.
Learning Rate Comparison	Hyperparameter Study	Compare convergence behaviour for different learning rates.
Decision Boundary (Optional)	Visualization	Illustrate the separating hyperplane using two selected features.

All plots above are generated and interpreted in-line under their respective tasks in Section 6.

8. Performance Tables

Training Summary

Dataset Size	1372
Train/Test Split	1097 / 275 (80% / 20%)
Learning Rate	0.01
Epochs	50
End Weights	$[-0.1941, -0.2358, -0.2109, -0.0195]$
End Bias	-0.10
Accuracy	0.9855
Precision	0.9683
Recall	1.0000
F1-score	0.9839

Epoch-wise Learning (Learning Rate = 0.01)

Epoch	Errors	Weight 1 (Variance)	Weight 2 (Skewness)	Bias
1	59	-0.0680	-0.0938	-0.03
5	20	-0.0991	-0.1295	-0.05
10	20	-0.1307	-0.1541	-0.06
20	19	-0.1636	-0.1946	-0.07
30	17	-0.1751	-0.2043	-0.08
40	18	-0.1948	-0.2244	-0.09
45	18	-0.1864	-0.2389	-0.10
46	15	-0.1942	-0.2323	-0.09
47	16	-0.1987	-0.2279	-0.09
48	17	-0.1951	-0.2336	-0.10
49	14	-0.2092	-0.2231	-0.10
50	18	-0.1941	-0.2358	-0.10

9. Additional Tasks

9.1 Step vs. Sigmoid Activation

The Step Activation Function gives us a simple yes or no answer it is either zero or one. It does not tell us how to get to that answer. This is a problem because we use gradients to find the way to do things. The Step Activation Function does not give us any information about how to change the weights in a -layer network because its gradient is zero everywhere except at one point where we cannot even use it.

The Sigmoid function is really different from the Step Activation Function. It gives us an answer that we can use to figure out the probability of something. The Sigmoid function is also very good, for learning because it gives us a gradient that's not zero so we can use it to update the weights in a multi-layer network. This makes it a lot easier to train the network and get results. The Sigmoid function is a choice when we want to use gradient descent to train a network.

9.2 Comparison with Scikit-learn's Perceptron

Both the from-scratch implementation and `sklearn.linear_model.Perceptron` follow the same perceptron learning rule and achieved comparable performance (accuracy ≈ 0.985 on the test set), with minor differences attributable to implementation details such as sample shuffling order and internal tie-breaking during updates.

9.3 Learning Rates 0.001, 0.01, 0.1

The three learning rates all gave the result. They all had the convergence curve and the same test accuracy of 0.9855. When we use a single-layer perceptron with the step activation function the learning rates change how big each update is. They do not change the direction of the update. So the sequence of mistakes the perceptron makes at each epoch is the same. The learning rates only

change the size of the weights and bias. The learning rates are important for the perceptron to learn the weights and bias.

9.4 Why a Single Layer Perceptron Cannot Solve XOR

XOR is not linearly separable: no single straight line can separate $(0,0) \rightarrow 0$, $(1,1) \rightarrow 0$ from $(0,1) \rightarrow 1$, $(1,0) \rightarrow 1$. Since a single-layer perceptron can only represent a linear decision boundary $\mathbf{w}^T \mathbf{x} + b = 0$, it cannot represent the XOR function. Solving XOR requires at least one hidden layer (a Multilayer Perceptron) with non-linear activations to construct a non-linear decision region. This is demonstrated experimentally in Appendix A.

9.5 Effect of Feature Normalization on Convergence

Standardizing the four features before training makes sure all inputs have similar scales. This stops features with larger raw values, like Skewness, from overshadowing smaller-scale features like Entropy during weight updates. In this experiment, using normalized data for training led to a more stable results than if we had used raw feature values. Unnormalized inputs usually result in slower and more erratic convergence.

10. Discussion

The Single Layer Perceptron achieved a high test accuracy of 98.55 on the Banknote Authentication dataset. It had perfect recall (1.0000) for the forged class and misclassified only four authentic notes. The exploratory data analysis showed that Variance and Skewness are the most distinct features, with the highest correlation to the class label. This was visually supported by the scatter plot and the two-feature decision boundary. Although the training error did not drop to zero within 50 epochs, this indicates that the four-dimensional feature space is not perfectly linearly separable; the classifier still performed well on new data. The study on learning rate revealed that, for this step-activation perceptron, the choice of η only influenced the scale of the learned parameters and did not affect the classification result.

11. Conclusion

I built a Single Layer Perceptron from the ground up. This included initializing weights and biases, using the step activation function, performing forward propagation, and applying the perceptron learning rule. I trained and tested the model on the UCI Banknote Authentication dataset, which resulted in a test accuracy of 98.55. The experiment confirmed the theoretical limits of the perceptron, such as its inability to solve non-linearly separable problems like XOR. It also showed the model's effectiveness on a nearly linearly separable real-world dataset. Additionally, it underscored the need for preprocessing, specifically normalization, and the shift to differentiable activation functions in deeper models.

12. References

1. F. Rosenblatt, "The Perceptron," *Psychological Review*, 1958.
2. I. Goodfellow, Y. Bengio and A. Courville, *Deep Learning*, MIT Press, 2016.

3. C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006.
4. S. Haykin, *Neural Networks and Learning Machines*, Pearson, 2009.
5. UCI Machine Learning Repository – Banknote Authentication Dataset.
6. Scikit-learn Documentation: Perceptron.

A Appendix A: Perceptron Learning on Logic Gates

A.1 Objective

This appendix extends the Single Layer Perceptron implementation from Experiment 1 to the classical logic gates AND, OR, NOT, and XOR. For each gate, the perceptron is trained from zero-initialized weights and bias, the weights are logged after every update, and the decision boundary is plotted after each update. The gate truth tables (4 rows for AND/OR/XOR, 2 rows for NOT) serve directly as the training data – no external dataset is required, since these truth tables are exhaustive.

A.2 AND Gate

```
Initial weights: [0. 0.], bias: 0.0

Update 1 (Epoch 1): input=[0 0], target=0, pred=1, weights=[0. 0.], bias=-0.10
Update 2 (Epoch 1): input=[1 1], target=1, pred=0, weights=[0.1 0.1], bias=0.00
Update 3 (Epoch 2): input=[0 0], target=0, pred=1, weights=[0.1 0.1], bias=-0.10
Update 4 (Epoch 2): input=[0 1], target=0, pred=1, weights=[0.1 0. ], bias=-0.20
Update 5 (Epoch 2): input=[1 1], target=1, pred=0, weights=[0.2 0.1], bias=-0.10
Update 6 (Epoch 3): input=[0 1], target=0, pred=1, weights=[0.2 0. ], bias=-0.20
Update 7 (Epoch 3): input=[1 0], target=0, pred=1, weights=[0.1 0. ], bias=-0.30
Update 8 (Epoch 3): input=[1 1], target=1, pred=0, weights=[0.2 0.1], bias=-0.20

Converged after epoch 4, total updates: 8
```




Figure 11: AND Gate – decision boundary after each of the 8 weight updates (initial state through convergence)

A.3 OR Gate

Initial weights: [0. 0.], bias: 0.0

Update 1 (Epoch 1): input=[0 0], target=0, pred=1, weights=[0. 0.], bias=-0.10
 Update 2 (Epoch 1): input=[0 1], target=1, pred=0, weights=[0. 0.1], bias=0.00
 Update 3 (Epoch 2): input=[0 0], target=0, pred=1, weights=[0. 0.1], bias=-0.10
 Update 4 (Epoch 2): input=[1 0], target=1, pred=0, weights=[0.1 0.1], bias=0.00
 Update 5 (Epoch 3): input=[0 0], target=0, pred=1, weights=[0.1 0.1], bias=-0.10

Converged after epoch 4, total updates: 5




Figure 12: OR Gate – decision boundary after each of the 5 weight updates (initial state through convergence)

A.4 NOT Gate

Initial weights: [0.], bias: 0.0

Update 1 (Epoch 1): input=[1], target=0, pred=1, weights=[-0.1], bias=-0.10
 Update 2 (Epoch 2): input=[0], target=1, pred=0, weights=[-0.1], bias=0.00

Converged after epoch 3, total updates: 2




Figure 13: NOT Gate – decision boundary after each of the 2 weight updates (initial state through convergence)

A.5 XOR Gate (Single Layer Perceptron)

```
Initial weights: [0. 0.], bias: 0.0

Update 1 (Epoch 1): input=[0 0], target=0, pred=1, weights=[0. 0.], bias=-0.10
Update 2 (Epoch 1): input=[0 1], target=1, pred=0, weights=[0. 0.1], bias=0.00
Update 3 (Epoch 1): input=[1 1], target=0, pred=1, weights=[-0.1 0.], bias=-0.10
Update 4 (Epoch 2): input=[0 1], target=1, pred=0, weights=[-0.1 0.1], bias=0.00
Update 5 (Epoch 2): input=[1 0], target=1, pred=0, weights=[0. 0.1], bias=0.10
Update 6 (Epoch 2): input=[1 1], target=0, pred=1, weights=[-0.1 0.], bias=0.00
Update 7 (Epoch 3): input=[0 0], target=0, pred=1, weights=[-0.1 0.], bias=-0.10
Update 8 (Epoch 3): input=[0 1], target=1, pred=0, weights=[-0.1 0.1], bias=0.00
Update 9 (Epoch 3): input=[1 0], target=1, pred=0, weights=[0. 0.1], bias=0.10
Update 10 (Epoch 3): input=[1 1], target=0, pred=1, weights=[-0.1 0.], bias=0.00

... From Epoch 3 onward, the perceptron repeats this exact same
4-update cycle (updates 7-10 above) identically every single
epoch, all the way through Epoch 20 ...

Update 75 (Epoch 20): input=[0 0], target=0, pred=1, weights=[-0.1 0.], bias=-0.10
Update 76 (Epoch 20): input=[0 1], target=1, pred=0, weights=[-0.1 0.1], bias=0.00
Update 77 (Epoch 20): input=[1 0], target=1, pred=0, weights=[0. 0.1], bias=0.10
Update 78 (Epoch 20): input=[1 1], target=0, pred=1, weights=[-0.1 0.], bias=0.00

Did NOT converge within 20 epochs, total updates: 78
```




Figure 14: XOR Gate – decision boundary after each weight update (first 12 states shown). The boundary line keeps swinging between two positions and never settles, visually confirming non-convergence.

A.6 Conclusion: The Single Layer Perceptron Does Not Suit XOR

The findings shown above prove that a Single Layer Perceptron **is not suitable for the XOR gate**. Whereas with AND, OR, and NOT — each of which converged after just a few updates — the XOR perceptron was unable to converge even after 20 complete epochs (78 weight updates), since from the third epoch it began to cycle indefinitely between the same four corrective updates and never attained a state in which there were zero misclassifications.

This occurs since XOR is not linearly separable; the positive class $\{(0, 1), (1, 0)\}$ and the negative class $\{(0, 0), (1, 1)\}$ are situated on opposite diagonals and therefore cannot be separated by any single straight line ($\mathbf{w}^T \mathbf{x} + b = 0$). As a single-layer perceptron can only ever correspond to one straight decision boundary, it is fundamentally unable to represent the XOR function — a fact that is visually evident from the oscillating boundary line in the plots mentioned above. In order to solve XOR it is necessary to have at least one hidden layer (that is, to use a Multilayer Perceptron) so as to combine several linear boundaries and form a non-linear decision region; this extension, however, is beyond the scope of the present single-layer perceptron appendix.